

Srikar Vadlamani

(412) 592-1238 | ssvadlam@andrew.cmu.edu | [linkedin.com/in/srikarv1](https://www.linkedin.com/in/srikarv1) | github.com/srikarv1 | [Google Scholar](#) | [Personal Website](#)

EDUCATION

Carnegie Mellon University

Pittsburgh, PA

Master of Information Systems Management

August 2025 – December 2026

Courses: AI Agents, Distributed Systems, Deep Learning, Machine Learning in Production, OOPS

GPA: 3.95

Delhi Technological University

Delhi, India

Bachelor of Technology in Computer Engineering | Minor in Machine Learning

August 2019 – May 2023

Courses: Advanced Data Structures, Operating Systems, Database Systems, Compiler Design

SKILLS

Programming Languages: Java, Python, C/C++, SQL Server, SQLite, Typescript, Javascript, Node.js

AI/ML: PyTorch, pandas, HuggingFace, TensorFlow, FastMCP, Model Context Protocol (MCP), Cursor, Windsurf, LLM agents

Cloud and Tools: AWS, GCP, Redis, Spark, Git, MongoDB, MySQL, Docker, Kubernetes, Vercel

EXPERIENCE

Adobe Inc. [CMU Capstone]

Pittsburgh, PA

Graduate Researcher

August 2026 - Present

- Spearheaded end-to-end **post-training of an RL based deep research orchestrator** using behavior cloning, DPO on counterfactual pairs, and GRPO on a LoRA adapted Qwen3-1.7B

Amazon Web Services Inc. - Redshift Query Optimization

East Palo Alto, CA

Software Development Engineer Intern

June 2026 – August 2026

- Architected and launched **Query Level Insights**, a Redshift diagnostics platform that lets engineers trace regressions from benchmark run to SQL statement, plan node, segment, and slice using plan trees and execution timelines
- Reduced core comparison scans from approximately 120 to 6 for a typical 20-template workload through grouped Redshift queries, materialized-view fallback, and S3 caching; backed the feature with 259 automated tests

The Bank of New York Mellon

Chennai, India

Full Stack Engineer II, Markets

July 2023 – May 2025

- Containerized and orchestrated **Java SpringBoot** microservices using **Docker**, **Kubernetes**, and **Helm**; managed multi-region FX hedging workflows and currency risk for 200+ global institutions
- Re-engineered a trade pipeline by integrating an **Apache Kafka** and **JMS** messaging backbone (10K+ trades/day) with **Hazelcast** distributed caching, slashing **MySQL** load by 70% and trade processing latency by 30%
- Owned the full Production Lifecycle by automating build-to-deploy workflows via **GitLab CI/CD** (50% faster releases); provided Level-3 Support for high-priority incidents utilizing **Grafana** and **Splunk**

Software Engineering Intern

May 2022 – July 2022

- Built **REST APIs** to ingest over **5M+ FX quotes/day** from third party sources into BNY's data warehouse, powering an AI-driven analytics platform used by investment bankers for trading insights
- Ensured scalable, fault-tolerant architectural systems by achieving **95%+ unit test coverage** using Mockito/JUnit and collaborating with senior engineers in **design and code reviews**

Samsung Innovation Lab, DTU

Delhi, India

Machine Learning Research Assistant

August 2022 – December 2022

- Reduced misclassification of steps on Samsung wearables by developing a **compressed 1D-CNN model** for gait analysis using **TensorFlow**; achieved 92% classification accuracy while reducing model size by 60% for on device inference

PROJECTS

Movie Recommendation Engine | Tech: Python, Kafka, SQLite, Flask, Prometheus, Grafana [CMU-17645]

- Engineered a movie recommendation system processing 223K+ real-time Kafka events via a **matrix factorization** model with SQLite persistence and sub-100ms Flask REST API, backed by an evaluation framework tracking **Precision@K**, **Click-Through Rate**, and **Match Rate** through real-time telemetry matching across **16K+ users**